

prompt_testing

February 16, 2026

```
[1]: # %pip install seaborn matplotlib scikit-learn --break-system-packages  
# %pip install tqdm ipywidgets notebook
```

```
[2]: import os  
import sys  
import json  
import time  
import glob  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from tqdm.notebook import tqdm  
from dotenv import load_dotenv  
from sklearn.metrics import accuracy_score, confusion_matrix,  
    classification_report  
try:  
    from IPython.display import display  
except:  
    pass  
  
# Add src to path  
sys.path.append(os.path.abspath("src"))  
  
from src.client import AzureLLMClient  
from src.config import EvalConfig  
from src.utils import encode_image  
  
# Load environment variables  
load_dotenv()  
  
# Configuration  
CONFIG_PATH = "config/eval.yaml"  
DATA_PATH = "data/annotations_balanced.csv" # Or annotations.csv  
PROMPTS_DIR = "prompts/"  
  
# Load config
```

```

config = EvalConfig(CONFIG_PATH)

print("Imports complete and config loaded.")

```

Imports complete and config loaded.

0.1 2. Load and Sample Data

We select exactly 24 images, ensuring 2 images for each of the 12 hair types (1A-4C).

```

[3]: def get_balanced_subset(csv_path, samples_per_class=2):
    df = pd.read_csv(csv_path)
    # Ensure allowed types only
    allowed_types = [f"{n}{l}" for n in "1234" for l in "ABC"]
    df = df[df['type'].isin(allowed_types)]

    # Sample
    subset = df.groupby('type', group_keys=False).apply(lambda x: x.
        sample(min(len(x), samples_per_class), random_state=42))

    # Verify count
    print(f"Selected {len(subset)} images.")
    print(subset['type'].value_counts())
    return subset

test_df = get_balanced_subset(DATA_PATH, samples_per_class=2)
# Display first few rows
test_df.head()

```

Selected 24 images.

type	count
1A	2
1B	2
1C	2
2A	2
2B	2
2C	2
3A	2
3B	2
3C	2
4A	2
4B	2
4C	2

Name: count, dtype: int64

/tmp/ipykernel_45783/467219597.py:8: FutureWarning: DataFrameGroupBy.apply
operated on the grouping columns. This behavior is deprecated, and in a future
version of pandas the grouping columns will be excluded from the operation.
Either pass `include_groups=False` to exclude the groupings or explicitly select

```
the grouping columns after groupby to silence this warning.
subset = df.groupby('type', group_keys=False).apply(lambda x:
x.sample(min(len(x), samples_per_class), random_state=42))
```

```
[3]:          image_path type
0   data/Straight_1a_11.jpg    1A
17  data/Straight_1a_6.jpg    1A
20  data/Straight_1B_19.jpg    1B
37  data/Straight_1b_5.jpg    1B
40  data/Straight_1C_11.jpg    1C
```

0.2 3. Initialize Client and Classifier

We allow for a custom “Top-2” capable classifier. Note that standard prompts might only return one class, in which case Top-2 Accuracy = Top-1 Accuracy. To enable true Top-2 measurement, your prompt should request a JSON output with a `candidates` list, e.g., `{"hair_type": "4C", "candidates": ["4C", "4B"]}`.

```
[4]: class AdvancedHairClassifier:
    def __init__(self, client, config, prompt_text):
        self.client = client
        self.cfg = config # Use EvalConfig object
        self.prompt = prompt_text
        # Use allowed types from config or fallback
        self.allowed = set(config.evaluation.get("allowed_types", [f"{{n}}{l}" for n in "1234" for l in "ABC"]))

    def predict(self, image_b64):
        # Using settings from config object (EvalConfig.api is a dict)
        api_cfg = self.cfg.api

        exceptions = []
        for attempt in range(api_cfg.get('retries', 2) + 1):
            try:
                raw_response = self.client.classify(
                    model=api_cfg['model'],
                    system_prompt=self.prompt,
                    image_b64=image_b64,
                    temperature=api_cfg['temperature'],
                    timeout=api_cfg['timeout']
                )

                # Parse JSON
                try:
                    data = json.loads(raw_response)
                    primary = data.get("hair_type", "Unknown")
                except json.JSONDecodeError:
                    primary = "Unknown"

                # Try to find candidates list for Top-2
                if "candidates" in data:
                    return data["candidates"]
            except Exception as e:
                exceptions.append(str(e))
                if attempt < api_cfg.get('retries', 2):
                    print(f"Attempt {attempt+1} failed: {e}. Retrying...")
```

```

# Expecting format: {"candidates": ["1A", "1B"]} or similar
candidates = data.get("candidates", [])
if not candidates:
    # Fallback: if only hair_type is present, candidates is just [primary]
    candidates = [primary]

# Ensure primary is in candidates if not present
if primary not in candidates:
    candidates.insert(0, primary)

return {
    "primary": primary if primary in self.allowed else "Unknown",
    "candidates": [c for c in candidates if c in self.allowed],
    "raw": raw_response
}
except json.JSONDecodeError:
    return {"primary": "Unknown", "candidates": [], "raw": raw_response}

except Exception as e:
    exceptions.append(str(e))
    if attempt < api_cfg.get('retries', 2):
        time.sleep(1)

return {"primary": "Unknown", "candidates": [], "error": str(exceptions)}

# Initialize Client
api_key = os.getenv("AZUREOPENAI_API_KEY")
# key name fixed based on .env content
endpoint = os.getenv("AZUREOPENAI_API_ENDPOINT")
api_version = os.getenv("AZUREOPENAI_API_VERSION", "2024-02-15-preview")

if not api_key:
    print("WARNING: AZUREOPENAI_API_KEY not found in environment. Check your .env file.")
if not endpoint:
    print("WARNING: AZUREOPENAI_API_ENDPOINT not found in environment. Check your .env file.")

# If endpoint is None, AzureOpenAI might fail or use default if configured elsewhere.

```

```

client = AzureLLMClient(api_key, endpoint, api_version)
print("Client initialized.")

```

Client initialized.

0.3 4. Run Evaluation on All Prompts

This loop iterates through all .txt files in the prompts/ directory.

```

[6]: # Define neighbor groups based on visual similarity (from run.ipynb)
NEIGHBORS = {
    "1A": ["1A", "1B"],
    "1B": ["1A", "1B", "1C"],
    "1C": ["1B", "1C", "2A"],
    "2A": ["1C", "2A", "2B"],
    "2B": ["2A", "2B", "2C"],
    "2C": ["2B", "2C", "3A"],
    "3A": ["2C", "3A", "3B"],
    "3B": ["3A", "3B", "3C"],
    "3C": ["3B", "3C", "4A"],
    "4A": ["3C", "4A", "4B"],
    "4B": ["4A", "4B", "4C"],
    "4C": ["4B", "4C"]
}

# Provide a helper function to calculate Top-K and Neighbor accuracy
def calculate_metrics(results_df):
    total = len(results_df)
    if total == 0:
        return {}

    # Top-1: Primary prediction matches true type
    top1_correct = sum(results_df['type'] == results_df['pred_primary'])

    # Neighbor Accuracy: True type is in the neighbor list of the predicted type
    def check_neighbor(row):
        pred = row['pred_primary']
        true_val = row['type']
        # Check if true label is in the neighbors of the predicted label
        if pred in NEIGHBORS:
            return true_val in NEIGHBORS[pred]
        return pred == true_val # Fallback to exact match if pred not in map

    neighbor_correct = results_df.apply(check_neighbor, axis=1).sum()

    # Top-2 (Standard): True type is in top 2 candidates
    def check_top2(row):
        cands = row['pred_candidates']

```

```

    top2 = cands[:2]
    return row['type'] in top2

top2_correct = results_df.apply(check_top2, axis=1).sum()

return {
    "Accuracy (Top-1)": top1_correct / total,
    "Accuracy (Neighbor)": neighbor_correct / total,
    # "Accuracy (Top-2)": top2_correct / total, # Keeping generic Top-2 if needed, but per request focusing on neighbors
    "Avg Latency (s)": results_df['latency'].mean(),
    "Median Latency (s)": results_df['latency'].median(),
    "P95 Latency (s)": results_df['latency'].quantile(0.95)
}

prompt_files = glob.glob(os.path.join(PROMPTS_DIR, "*.txt"))
prompt_files.sort()

all_results = {} # Store df per prompt
prompt_metrics = []

for p_file in prompt_files:
    p_name = os.path.basename(p_file)
    print(f"Testing Prompt: {p_name}")

    with open(p_file, 'r') as f:
        prompt_text = f.read()

    classifier = AdvancedHairClassifier(client, config, prompt_text)

    results = []

    # Using tqdm for progress bar
    for _, row in tqdm(test_df.iterrows(), total=len(test_df), desc=p_name):
        img_path = row['image_path']
        true_type = row['type']

        # Read image
        try:
            b64_img = encode_image(img_path)
        except Exception as e:
            print(f"Error reading {img_path}: {e}")
            continue

        start_t = time.time()
        pred_res = classifier.predict(b64_img)
        lat = time.time() - start_t

```

```

        results.append({
            "image_path": img_path,
            "type": true_type,
            "pred_primary": pred_res["primary"],
            "pred_candidates": pred_res["candidates"],
            "latency": lat,
            "raw_output": pred_res.get("raw", "")
        })

    res_df = pd.DataFrame(results)
    all_results[p_name] = res_df

    metrics = calculate_metrics(res_df)
    metrics["Prompt"] = p_name
    prompt_metrics.append(metrics)

    print(f"  Top-1: {metrics['Accuracy (Top-1)']:.2%}")
    print(f"  Neighbor Acc: {metrics['Accuracy (Neighbor)']:.2%}")

metrics_df = pd.DataFrame(prompt_metrics).set_index("Prompt")
display(metrics_df)

```

Testing Prompt: vision_only.txt

vision_only.txt: 0% | 0/24 [00:00<?, ?it/s]

Top-1: 25.00%

Neighbor Acc: 95.83%

Testing Prompt: vision_only_2.txt

vision_only_2.txt: 0% | 0/24 [00:00<?, ?it/s]

Top-1: 29.17%

Neighbor Acc: 83.33%

Testing Prompt: vision_only_3.txt

vision_only_3.txt: 0% | 0/24 [00:00<?, ?it/s]

Top-1: 0.00%

Neighbor Acc: 0.00%

Testing Prompt: vision_only_4.txt

vision_only_4.txt: 0% | 0/24 [00:00<?, ?it/s]

Top-1: 0.00%

Neighbor Acc: 0.00%

Testing Prompt: vision_only_5.txt

vision_only_5.txt: 0% | 0/24 [00:00<?, ?it/s]

Top-1: 0.00%

Neighbor Acc: 0.00%

```

Testing Prompt: vision_only_6.txt
vision_only_6.txt: 0% | 0/24 [00:00<?, ?it/s]
    Top-1: 37.50%
    Neighbor Acc: 87.50%

          Accuracy (Top-1)  Accuracy (Neighbor)  Avg Latency (s)  \
Prompt
vision_only.txt        0.250000        0.958333      31.383989
vision_only_2.txt      0.291667        0.833333      32.646621
vision_only_3.txt      0.000000        0.000000      98.412900
vision_only_4.txt      0.000000        0.000000     138.172958
vision_only_5.txt      0.000000        0.000000     100.740393
vision_only_6.txt      0.375000        0.875000      31.043881

          Median Latency (s)  P95 Latency (s)
Prompt
vision_only.txt        4.040297       68.343331
vision_only_2.txt      31.758237      64.141274
vision_only_3.txt      105.164299      126.646988
vision_only_4.txt      104.136385      128.667849
vision_only_5.txt      119.851002      129.496880
vision_only_6.txt      30.845677       61.334328

```

0.4 5. Visual Reports & Misclassifications

Select a specific prompt to inspect (defaulting to the one with best Top-1 accuracy).

```
[7]: if not metrics_df.empty:
    best_prompt = metrics_df['Accuracy (Top-1)'].idxmax()
    print(f"Generating detailed report for Best Prompt: {best_prompt}")

    df_best = all_results[best_prompt]

    # 1. Confusion Matrix
    y_true = df_best['type']
    y_pred = df_best['pred_primary']
    labels = sorted(list(set(y_true) | set(y_pred)))

    cm = confusion_matrix(y_true, y_pred, labels=labels)

    plt.figure(figsize=(10, 8))
    sns.heatmap(cm, annot=True, fmt='d', xticklabels=labels,
                yticklabels=labels, cmap='Blues')
    plt.title(f'Confusion Matrix - {best_prompt}')
    plt.ylabel('True Label')
    plt.xlabel('Predicted Label')
    plt.show()
```

```

# 2. Latency Histogram
plt.figure(figsize=(10, 4))
sns.histplot(df_best['latency'], kde=True)
plt.title(f'Latency Distribution - {best_prompt}')
plt.xlabel('Seconds')
plt.show()

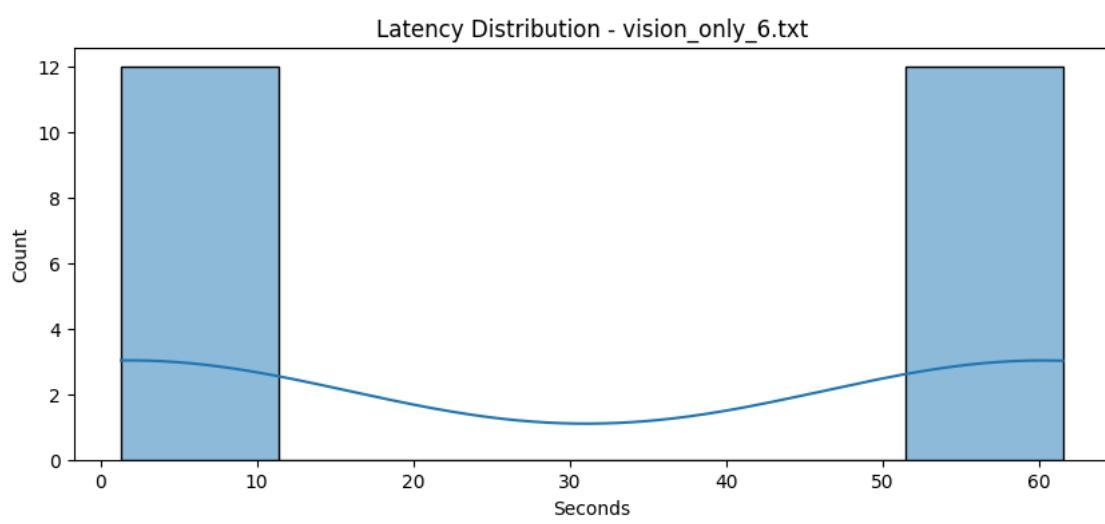
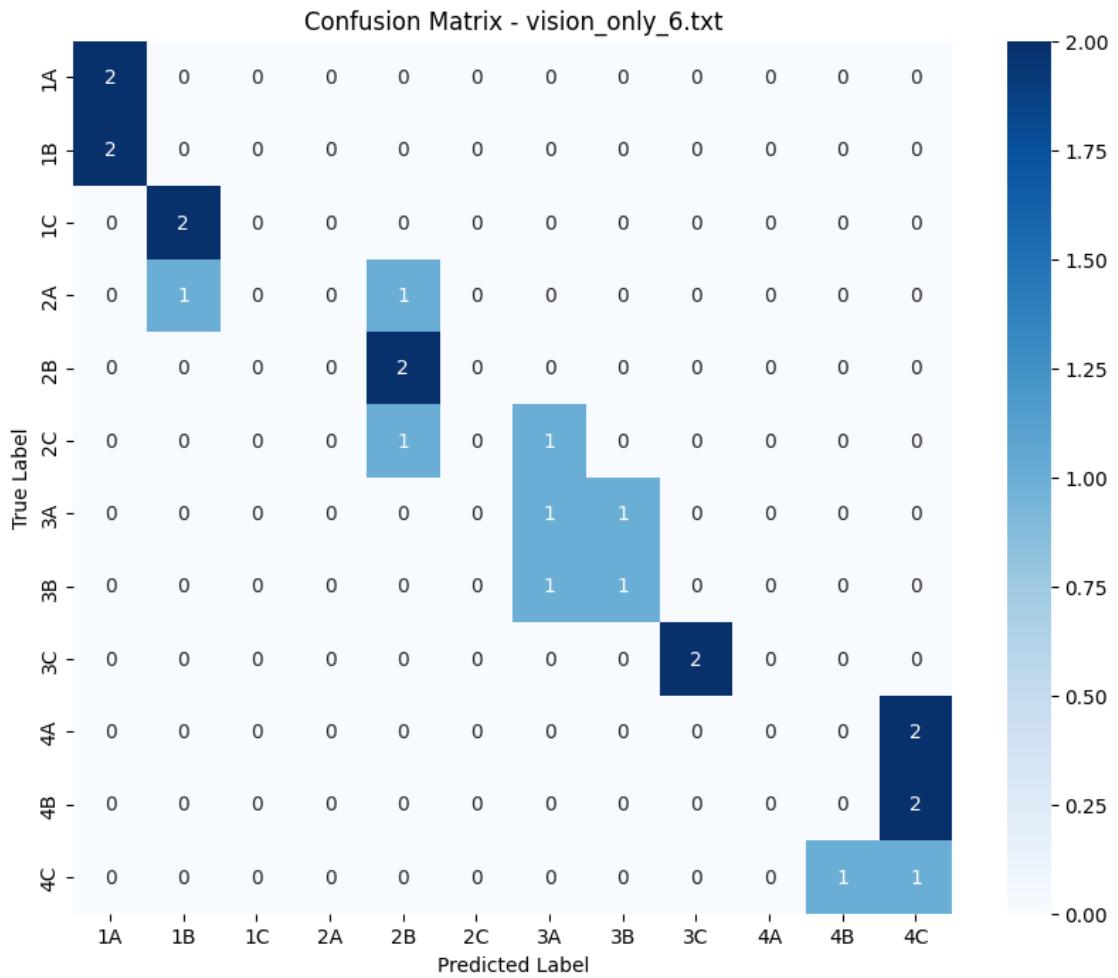
# 3. Misclassifications
misclassified = df_best[df_best['type'] != df_best['pred_primary']]
if not misclassified.empty:
    print(f"Found {len(misclassified)} misclassifications:")
    display(misclassified[['image_path', 'type', 'pred_primary', 'pred_candidates', 'raw_output']])

    # Visualize first few misclassifications
    # (Assuming we run this locally and can see images)
    print("Visualizing up to 3 misclassifications:")
    for idx, row in misclassified.head(3).iterrows():
        img_p = row['image_path']
        try:
            # If running in VS Code with local kernel, this works.
            # If paths are correct relative to cwd
            from PIL import Image
            img = Image.open(img_p)
            plt.figure(figsize=(3, 3))
            plt.imshow(img)
            plt.axis('off')
            plt.title(f"True: {row['type']} | Pred: {row['pred_primary']}")
            plt.show()
            print(f"Candidates: {row['pred_candidates']}")
            print("-" * 50)
        except Exception as e:
            print(f"Could not load image {img_p}: {e}")

else:
    print("No results to display.")

```

Generating detailed report for Best Prompt: vision_only_6.txt



Found 15 misclassifications:

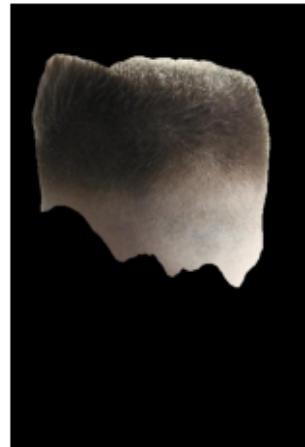
	image_path	type	pred_primary	pred_candidates	\
2	data/Straight_1B_19.jpg	1B	1A	[1A]	
3	data/Straight_1b_5.jpg	1B	1A	[1A]	
4	data/Straight_1C_11.jpg	1C	1B	[1B]	
5	data/Straight_1c_4.jpg	1C	1B	[1B]	
6	data/Wavy_2a_2.jpg	2A	2B	[2B]	
7	data/Wavy_2a_10.jpg	2A	1B	[1B]	
10	data/Wavy_2C_2.jpg	2C	3A	[3A]	
11	data/Wavy_2C_35.jpg	2C	2B	[2B]	
13	data/Curly_3A_9.jpg	3A	3B	[3B]	
15	data/new_3B_7.jpeg	3B	3A	[3A]	
18	data/Kinky_coily_4A_1.jpg	4A	4C	[4C]	
19	data/new_4A_1.jpeg	4A	4C	[4C]	
20	data/KInky_coily_4B_3.jpg	4B	4C	[4C]	
21	data/Kinky_coily_4B_18.jpg	4B	4C	[4C]	
23	data/new_4C_19.jpeg	4C	4B	[4B]	

raw_output

```
2 {\n  "hair_type": "1A"\n}\n3 {\n  "hair_type": "1A"\n}\n4 {\n  "hair_type": "1B"\n}\n5 {\n  "hair_type": "1B"\n}\n6 {\n  "hair_type": "2B"\n}\n7 {\n  "hair_type": "1B"\n}\n10 {\n  "hair_type": "3A"\n}\n11 {\n  "hair_type": "2B"\n}\n13 {\n  "hair_type": "3B"\n}\n15 {\n  "hair_type": "3A"\n}\n18 {\n  "hair_type": "4C"\n}\n19 {\n  "hair_type": "4C"\n}\n20 {\n  "hair_type": "4C"\n}\n21 {\n  "hair_type": "4C"\n}\n23 {\n  "hair_type": "4B"\n}
```

Visualizing up to 3 misclassifications:

True: 1B | Pred: 1A



Candidates: ['1A']

True: 1B | Pred: 1A



Candidates: ['1A']

True: 1C | Pred: 1B



Candidates: ['1B']

[]: